

CS2106 Operating Systems

Semester 1 2019/2020

Week 4 (02-06 September 2019)

Tutorial 2: Process Abstraction in Unix

1. (Behavior of `fork()` system call) The C program below attempts to highlight the behavior of the `fork()` system call:

C code:

```
int dataX = 100;
int main( )
{
    pid_t childPID;

    int dataY = 200;
    int* dataZptr = (int*) malloc(sizeof(int));

    *dataZptr = 300;

    //First Phase
    printf("PID[%d] | X = %d | Y = %d | Z = %d |\n",
           getpid(), dataX, dataY, *dataZptr);

    //Second Phase
    childPID = fork();
    printf("**PID[%d] | X = %d | Y = %d | Z = %d |\n",
           getpid(), dataX, dataY, *dataZptr);

    dataX += 1;
    dataY += 2;
    (*dataZptr) += 3;
    printf("#PID[%d] | X = %d | Y = %d | Z = %d |\n",
           getpid(), dataX, dataY, *dataZptr);

    //Insertion Point

    //Third Phase
    childPID = fork();
    printf("**PID[%d] | X = %d | Y = %d | Z = %d |\n",
           getpid(), dataX, dataY, *dataZptr);

    dataX += 1;
    dataY += 2;
    (*dataZptr) += 3;
    printf("##PID[%d] | X = %d | Y = %d | Z = %d |\n",
           getpid(), dataX, dataY, *dataZptr);

    return 0;
}
```

Please run the given program "**ForkTest.c**" on your system before answering the questions below.

- a. What is the difference between the 3 variables: **dataX**, **dataY**, **dataZptr**, and the memory location pointed to by **dataZptr**?
- b. Explain the **values** that are printed by the program.
- c. Focusing on the messages generated by second phase (they are prefixed with either "*" and "#"), what can you say about the behavior of the **fork()** system call?
- d. Using the messages seen on your system, draw a **process tree** to represent the processes generated. Use the process tree to explain the values printed by the child processes.
- e. Do you think it is possible to get different ordering between the output messages, why?
- f. Can you point out which pair(s) of messages can never swap places? i.e. their relative order is always the same?
- g. If we insert the following code at the insertion point:

Sleep Code
<pre>if (childPID == 0){ sleep(5); //sleep for 5 seconds }</pre>

How does this change the ordering of the output messages? State your assumption, if any.

- h. Instead of the code in (g), we insert the following code at the insertion point:

Wait Code
<pre>if (childPID != 0){ wait(NULL); //NULL means we don't care // about the return result }</pre>

How does this change the ordering of the output messages? State your assumption, if any.

2. (Process Creation) The following program calculates factorial of a given number. The source code **FF.c** is also given for your own test.

C code:
<pre>int factorial(int n) { if (n == 0){ fork(); // NOTE the change return 1; } return factorial(n-1) * n; } int main() { printf("fac(2) = %d\n", factorial(2)); return 0; }</pre>

- Give and explain the execution output.
- If the line of **fork()** is moved above the **if** statement, what is the execution output?

New Factorial Code
<pre>int factorial(int n) { fork(); // NOTE the change if (n == 0){ return 1; } return factorial(n-1) * n; }</pre>

- (Continue from b, source code in **FF_2.c**) How many lines of results will be printed if **fac(n)** is called instead?

New Main Code
<pre>int main() { int n; printf("Input n: "); scanf("%d", &n); printf("fac(%d) = %d\n", n, factorial(n)); return 0; }</pre>

3. (Parallel computation) Even with the crude synchronization mechanism, we can solve programming problems in new (and exciting) ways. We will attempt to utilize multiple processes to work on a problem simultaneously in this question.

You are given two C source code "**Parallel.c**" and "**PrimeFactors.c**". The "**PrimeFactors.c**" is a simple prime factorization program. "**Parallel.c**" use the "**fork()**" and "**execl()**" combination to spawn off a new process to run the prime factorization.

Let's setup the programs as follows:

1. Compile "**PrimeFactors.c**" to get a executable with name "**PF**":
gcc PrimeFactors.c -o PF
2. Compiles "**Parallel.c**": **gcc Parallel.c**

Run the **a.out** generated from step (2). Below is a sample session:

```
$> a.out
1024
1024 has 10 prime factors //note: not unique prime factors
```

If you try large prime numbers, e.g. 111113111, the program may take a while.

Modify only Parallel.c such that we can now initiate prime factorization on [1-9] user inputs simultaneously. More importantly, we want to report result as soon as they are ready regardless of the user input order.

Sample session below:

```
$> a.out < test2.in
9 has 2 prime factors //Results
118689518 has 3 prime factors
44721359 has 1 prime factors
99999989 has 1 prime factors
111113111 has 1 prime factors
```

Note the order of the result may differ on your system. Most of time, they should follow roughly the computation time needed (composite number < prime number and small number < large number). Two simple test cases are given "**test1.in**" and "**test2.in**" to aid your testing.

Most of what you need is already demonstrated in the original **Parallel.c** (so that this is more of a mechanism question rather than a coding question). You only need "**fork()**", "**execl()**" and "**wait()**" for your solution.

After you have solved the problem, find a way to change your **wait()** to **waitpid()**, **what do you think is the effect of this change?**

Additional Questions (For exploration only, not discussed in tutorial)

1. (Process Creation) Consider the following sequence of instructions in a C program:

C code:
<pre>int x = 10; int y = 123; y = fork(); if (y == 0) x--; y = fork(); if (y == 0) x--; printf("[PID %d]: x=%d, y=%d\n", getpid(), x, y);</pre>

You can assume that the first process has process number 100 (and so `getpid()` returns the value 100 for this process), and that the processes created (in order) are 101,102 and so on.

Give:

- A possible final set of printed messages.
- An impossible final set of printed messages.